

Inam_Daudi_CIND820

November 9, 2021

```
[ ]: #import libraries
      #import pandas as pd
      #import numpy as np
      #import matplotlib.pyplot as plt
      #from matplotlib.patches import Patch
      #from matplotlib.lines import Line2D
      #import seaborn as sns
      #import sklearn as sk
      #from sklearn import metrics
      #from sklearn.metrics import classification_report
      #from sklearn.linear_model import LogisticRegression
      #from sklearn.tree import DecisionTreeClassifier
      #from sklearn.neighbors import KNeighborsClassifier
      #from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
      #from sklearn.naive_bayes import GaussianNB
      #from sklearn.svm import SVC
      #from sklearn.metrics import confusion_matrix
      #from sklearn.ensemble import BaggingClassifier
      #from sklearn.ensemble import GradientBoostingClassifier
```

```
[144]: #import data
        coupon_data = pd.read_csv('in-vehicle-coupon-recommendation.csv')
```

```
[145]: #data attributes and their types
        coupon_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12684 entries, 0 to 12683
Data columns (total 26 columns):
 #   Column              Non-Null Count  Dtype
---  -
 0   destination         12684 non-null  object
 1   passenger           12684 non-null  object
 2   weather             12684 non-null  object
 3   temperature         12684 non-null  int64
 4   time                12684 non-null  object
 5   coupon              12684 non-null  object
 6   expiration          12684 non-null  object
```

```

7   gender                12684 non-null  object
8   age                   12684 non-null  object
9   maritalStatus         12684 non-null  object
10  has_children           12684 non-null  int64
11  education              12684 non-null  object
12  occupation             12684 non-null  object
13  income                 12684 non-null  object
14  car                    108 non-null   object
15  Bar                    12577 non-null  object
16  CoffeeHouse           12467 non-null  object
17  CarryAway              12533 non-null  object
18  RestaurantLessThan20  12554 non-null  object
19  Restaurant20To50      12495 non-null  object
20  toCoupon_GEQ5min       12684 non-null  int64
21  toCoupon_GEQ15min      12684 non-null  int64
22  toCoupon_GEQ25min      12684 non-null  int64
23  direction_same        12684 non-null  int64
24  direction_opp         12684 non-null  int64
25  Y                      12684 non-null  int64

```

dtypes: int64(8), object(18)

memory usage: 2.5+ MB

```
[146]: #first few tuples in the dataset
coupon_data.head()
```

```
[146]:
      destination  passanger  weather  temperature  time \
0  No Urgent Place    Alone   Sunny           55   2PM
1  No Urgent Place  Friend(s)   Sunny           80  10AM
2  No Urgent Place  Friend(s)   Sunny           80  10AM
3  No Urgent Place  Friend(s)   Sunny           80   2PM
4  No Urgent Place  Friend(s)   Sunny           80   2PM

```

```

      coupon expiration  gender  age  maritalStatus  ... \
0  Restaurant(<20)      1d  Female  21  Unmarried partner  ...
1  Coffee House        2h  Female  21  Unmarried partner  ...
2  Carry out & Take away  2h  Female  21  Unmarried partner  ...
3  Coffee House        2h  Female  21  Unmarried partner  ...
4  Coffee House        1d  Female  21  Unmarried partner  ...

```

```

      CoffeeHouse  CarryAway  RestaurantLessThan20  Restaurant20To50  \
0      never      NaN      4~8      1~3
1      never      NaN      4~8      1~3
2      never      NaN      4~8      1~3
3      never      NaN      4~8      1~3
4      never      NaN      4~8      1~3

```

```

toCoupon_GEQ5min  toCoupon_GEQ15min  toCoupon_GEQ25min  direction_same  \

```

0	1	0	0	0
1	1	0	0	0
2	1	1	0	0
3	1	1	0	0
4	1	1	0	0

	direction_opp	Y
0	1	1
1	1	0
2	1	1
3	1	0
4	1	0

[5 rows x 26 columns]

```
[147]: #data description
coupon_data.describe(include='all')
```

```
[147]:
```

	destination	passanger	weather	temperature	time	coupon \
count	12684	12684	12684	12684.000000	12684	12684
unique	3	4	3	NaN	5	5
top	No Urgent Place	Alone	Sunny	NaN	6PM	Coffee House
freq	6283	7305	10069	NaN	3230	3996
mean	NaN	NaN	NaN	63.301798	NaN	NaN
std	NaN	NaN	NaN	19.154486	NaN	NaN
min	NaN	NaN	NaN	30.000000	NaN	NaN
25%	NaN	NaN	NaN	55.000000	NaN	NaN
50%	NaN	NaN	NaN	80.000000	NaN	NaN
75%	NaN	NaN	NaN	80.000000	NaN	NaN
max	NaN	NaN	NaN	80.000000	NaN	NaN

	expiration	gender	age	maritalStatus	...	CoffeeHouse	CarryAway \
count	12684	12684	12684	12684	...	12467	12533
unique	2	2	8	5	...	5	5
top	1d	Female	21	Married partner	...	less1	1~3
freq	7091	6511	2653	5100	...	3385	4672
mean	NaN	NaN	NaN	NaN	...	NaN	NaN
std	NaN	NaN	NaN	NaN	...	NaN	NaN
min	NaN	NaN	NaN	NaN	...	NaN	NaN
25%	NaN	NaN	NaN	NaN	...	NaN	NaN
50%	NaN	NaN	NaN	NaN	...	NaN	NaN
75%	NaN	NaN	NaN	NaN	...	NaN	NaN
max	NaN	NaN	NaN	NaN	...	NaN	NaN

	RestaurantLessThan20	Restaurant20To50	toCoupon_GEQ5min \
count	12554	12495	12684.0
unique	5	5	NaN

top	1~3	less1	NaN
freq	5376	6077	NaN
mean	NaN	NaN	1.0
std	NaN	NaN	0.0
min	NaN	NaN	1.0
25%	NaN	NaN	1.0
50%	NaN	NaN	1.0
75%	NaN	NaN	1.0
max	NaN	NaN	1.0

	toCoupon_GEQ15min	toCoupon_GEQ25min	direction_same	direction_opp \
count	12684.000000	12684.000000	12684.000000	12684.000000
unique	NaN	NaN	NaN	NaN
top	NaN	NaN	NaN	NaN
freq	NaN	NaN	NaN	NaN
mean	0.561495	0.119126	0.214759	0.785241
std	0.496224	0.323950	0.410671	0.410671
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	1.000000
50%	1.000000	0.000000	0.000000	1.000000
75%	1.000000	0.000000	0.000000	1.000000
max	1.000000	1.000000	1.000000	1.000000

	Y
count	12684.000000
unique	NaN
top	NaN
freq	NaN
mean	0.568433
std	0.495314
min	0.000000
25%	0.000000
50%	1.000000
75%	1.000000
max	1.000000

[11 rows x 26 columns]

```
[148]: #change the temperature attribute to the 'category' type
coupon_data['temperature']=coupon_data['temperature'].astype('category')
```

```
[149]: #checking the types and counts of the cars attribute
coupon_data["car"].value_counts()
```

```
[149]: Mazda5                22
       Scooter and motorcycle  22
       do not drive          22
```

```

Car that is too old to install Onstar :D      21
crossover                                     21
Name: car, dtype: int64

```

```

[150]: #dropping the car column
coupon_data.drop('car', inplace=True, axis=1)

```

```

[151]: #fill missing values with the most common value
for x in coupon_data.columns[coupon_data.isna().any()]:
    coupon_data = coupon_data.fillna({x: coupon_data[x].value_counts().
    ↳idxmax()})

```

```

[154]: #changing object datatypes to categorical datatypes
coupon_data_obj = coupon_data.select_dtypes(include=['object']).copy()

for col in coupon_data_obj.columns:
    coupon_data[col]=coupon_data[col].astype('category')

```

```

[155]: #counting unique values in the columns with numeric values
coupon_data.select_dtypes('int64').nunique()

```

```

[155]: has_children      2
toCoupon_GEQ5min        1
toCoupon_GEQ15min       2
toCoupon_GEQ25min       2
direction_same          2
direction_opp           2
Y                       2
dtype: int64

```

```

[156]: #dropping the toCoupon_GEQ5min column as it only has 1 value
coupon_data.drop(columns=['toCoupon_GEQ5min'], inplace=True)

```

```

[157]: #data attributes and their types
coupon_data.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12684 entries, 0 to 12683
Data columns (total 24 columns):
#   Column                Non-Null Count  Dtype
---  -
0   destination            12684 non-null  category
1   passanger              12684 non-null  category
2   weather                12684 non-null  category
3   temperature            12684 non-null  category
4   time                   12684 non-null  category
5   coupon                 12684 non-null  category

```

6	expiration	12684	non-null	category
7	gender	12684	non-null	category
8	age	12684	non-null	category
9	maritalStatus	12684	non-null	category
10	has_children	12684	non-null	int64
11	education	12684	non-null	category
12	occupation	12684	non-null	category
13	income	12684	non-null	category
14	Bar	12684	non-null	category
15	CoffeeHouse	12684	non-null	category
16	CarryAway	12684	non-null	category
17	RestaurantLessThan20	12684	non-null	category
18	Restaurant20To50	12684	non-null	category
19	toCoupon_GEQ15min	12684	non-null	int64
20	toCoupon_GEQ25min	12684	non-null	int64
21	direction_same	12684	non-null	int64
22	direction_opp	12684	non-null	int64
23	Y	12684	non-null	int64

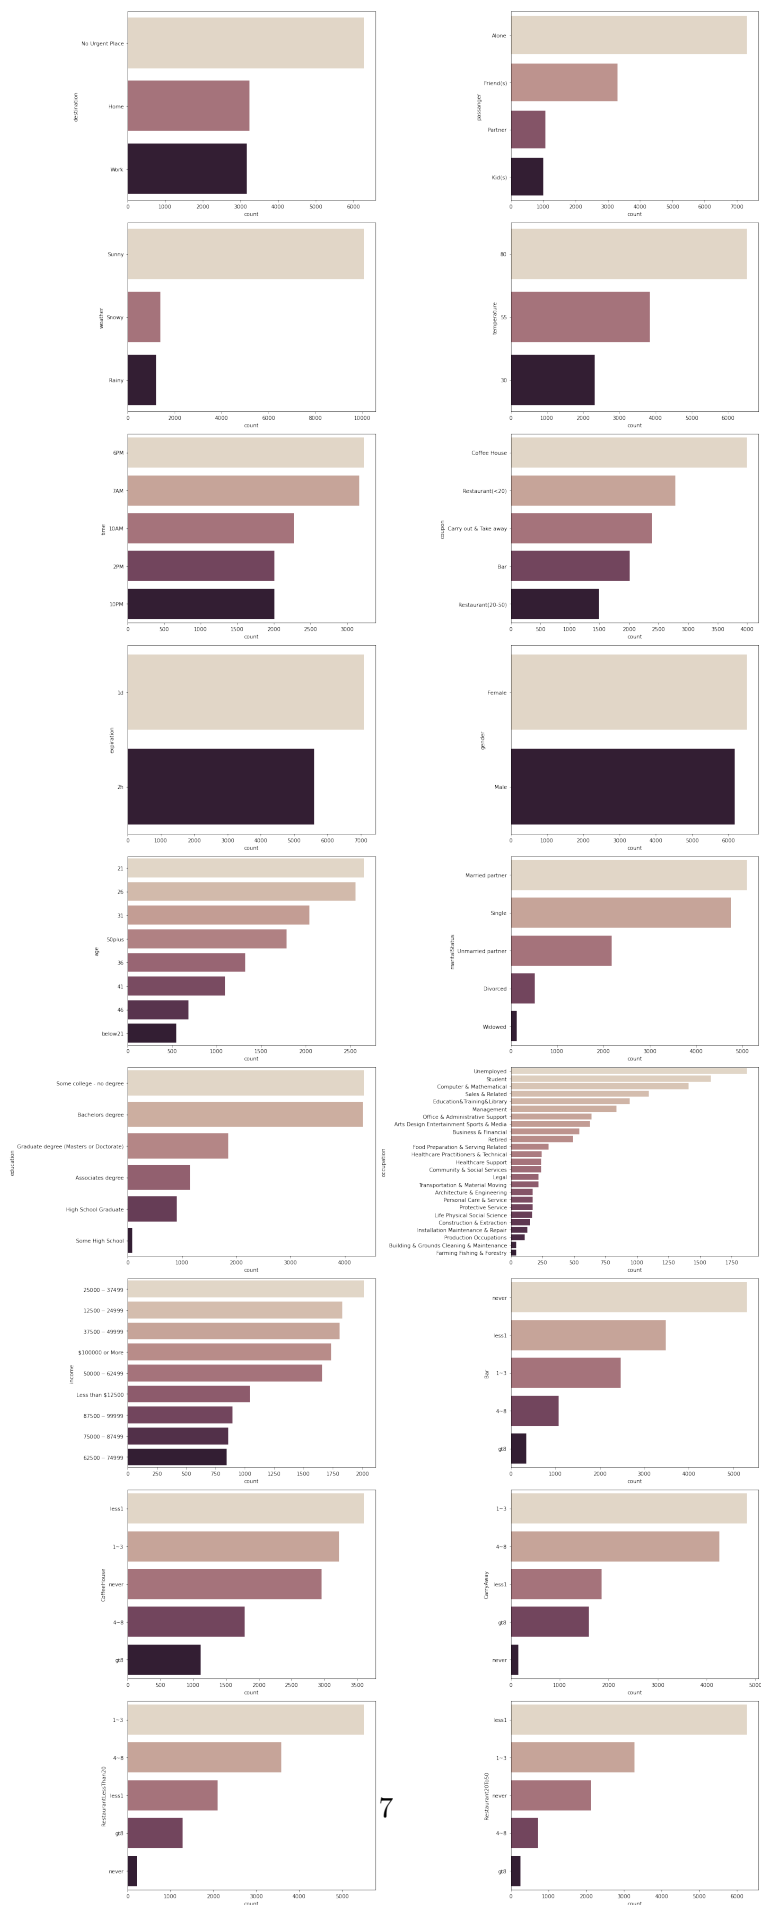
dtypes: category(18), int64(6)

memory usage: 821.7 KB

```
[158]: #plotting bar charts for all categorical attributes
fig, axes = plt.subplots(9, 2, figsize=(20,50))
axes = axes.flatten()

for ax, col in zip(axes, coupon_data.select_dtypes('category').columns):
    sns.countplot(y=col, data=coupon_data, ax=ax,
                  palette="ch:.25", order=coupon_data[col].value_counts().
    ↪index);

plt.tight_layout()
plt.show()
```



```
[159]: #Transforming the data using LabelEncoder and OneHotEncoder
from sklearn.preprocessing import LabelEncoder, OneHotEncoder

enc = OneHotEncoder(dtype='int64')

coupon_data_cat = coupon_data.select_dtypes(include=['category']).copy()
coupon_data_int = coupon_data.select_dtypes(include=['int64']).copy()

coupon_data_enc = pd.DataFrame()
for col in coupon_data_cat.columns:
    enc_results = enc.fit_transform(coupon_data_cat[[col]])
    coupon_data0 = pd.DataFrame(enc_results.toarray(), columns=enc.categories_)
    coupon_data_enc = pd.concat([coupon_data_enc, coupon_data0], axis=1)

coupon_data_final = pd.concat([coupon_data_enc, coupon_data_int], axis=1)
```

```
[160]: coupon_data_final
```

```
[160]:      (Home,)  (No Urgent Place,)  (Work,)  (Alone,)  (Friend(s),)  \
0           0           1           0           1           0
1           0           1           0           0           1
2           0           1           0           0           1
3           0           1           0           0           1
4           0           1           0           0           1
...      ...      ...      ...      ...      ...
12679      1           0           0           0           0
12680      0           0           1           1           0
12681      0           0           1           1           0
12682      0           0           1           1           0
12683      0           0           1           1           0

      (Kid(s),)  (Partner,)  (Rainy,)  (Snowy,)  (Sunny,)  ...  (4~8,)  \
0           0           0           0           0           1  ...  0
1           0           0           0           0           1  ...  0
2           0           0           0           0           1  ...  0
3           0           0           0           0           1  ...  0
4           0           0           0           0           1  ...  0
...      ...      ...      ...      ...      ...      ...
12679      0           1           1           0           0  ...  0
12680      0           0           1           0           0  ...  0
12681      0           0           0           1           0  ...  0
12682      0           0           0           1           0  ...  0
12683      0           0           0           0           1  ...  0

      (gt8,)  (less1,)  (never,)  has_children  toCoupon_GEQ15min  \
```


0	0	0	0	1	0
1	0	0	0	1	0
2	0	0	0	1	1
3	0	0	0	1	1
4	0	0	0	1	1
...
12679	0	0	0	0	0
12680	0	0	0	0	0
12681	0	0	0	0	0
12682	0	0	0	0	1
12683	0	0	0	0	0

	toCoupon_GEQ25min	direction_same	direction_opp	Y
0	0	0	1	1
1	0	0	1	0
2	0	0	1	1
3	0	0	1	0
4	0	0	1	0
...
12679	0	1	0	1
12680	0	0	1	1
12681	0	1	0	0
12682	1	0	1	0
12683	0	1	0	0

[12684 rows x 111 columns]

```
[161]: #Splitting the data into test and training
from sklearn.model_selection import train_test_split

train_set, test_set = train_test_split(coupon_data_final, test_size=.2,
    random_state=42, shuffle=True, stratify=df_final['Y'])
```

```
[162]: #Creating the dependent and independent variables
X_train = train_set.iloc[:, :-1].values
X_test = test_set.iloc[:, :-1].values

y_train = train_set.iloc[:, -1].values
y_test = test_set.iloc[:, -1].values
```

```
[163]: #Applying Logistic Regression
LR = LogisticRegression(random_state=0, solver='lbfgs', multi_class='ovr').
    fit(X_train, y_train)

y_pred_LR = LR.predict(X_test)

print(classification_report(y_test, y_pred_LR))
```

```
#Logistic Regression Confusion Matrix
cm_LR = confusion_matrix(y_test,y_pred_LR)
pd.crosstab(y_test, y_pred_LR, rownames = ['Truth'], colnames =['Predicted'],
↪margins = True)
```

	precision	recall	f1-score	support
0	0.65	0.56	0.60	1095
1	0.70	0.77	0.73	1442
accuracy			0.68	2537
macro avg	0.68	0.67	0.67	2537
weighted avg	0.68	0.68	0.68	2537

```
[163]: Predicted    0     1   All
      Truth
0         615   480  1095
1         327  1115  1442
All        942  1595  2537
```

```
[164]: #Applying Decision Tree
DTC = DecisionTreeClassifier().fit(X_train, y_train)

y_pred_DTC = DTC.predict(X_test)

print(classification_report(y_test, y_pred_DTC))

#Decision Tree Confusion Matrix
cm_DTC = confusion_matrix(y_test,y_pred_DTC)
pd.crosstab(y_test, y_pred_DTC, rownames = ['Truth'], colnames =['Predicted'],
↪margins = True)
```

	precision	recall	f1-score	support
0	0.62	0.62	0.62	1095
1	0.71	0.71	0.71	1442
accuracy			0.67	2537
macro avg	0.67	0.67	0.67	2537
weighted avg	0.67	0.67	0.67	2537

```
[164]: Predicted    0     1   All
      Truth
0         680   415  1095
```

1	412	1030	1442
All	1092	1445	2537

```
[165]: #Applying Decision Tree Ensemble Bagging
BC = BaggingClassifier().fit(X_train, y_train)

y_pred_BC = BC.predict(X_test)

print(classification_report(y_test, y_pred_BC))

#Decision Tree Ensemble Bagging Confusion Matrix
cm_BC = confusion_matrix(y_test,y_pred_BC)
pd.crosstab(y_test, y_pred_BC, rownames = ['Truth'], colnames =['Predicted'],
↪margins = True)
```

	precision	recall	f1-score	support
0	0.68	0.71	0.69	1095
1	0.77	0.75	0.76	1442
accuracy			0.73	2537
macro avg	0.73	0.73	0.73	2537
weighted avg	0.73	0.73	0.73	2537

```
[165]: Predicted    0    1   All
Truth
0          776   319 1095
1          364 1078 1442
All        1140 1397 2537
```

```
[166]: #Applying Decision Tree Ensemble Boosting
GBC = GradientBoostingClassifier().fit(X_train, y_train)

y_pred_GBC = GBC.predict(X_test)

print(classification_report(y_test, y_pred_GBC))

#Decision Tree Ensemble Gradient Boosting Confusion Matrix
cm_GBC = confusion_matrix(y_test,y_pred_GBC)
pd.crosstab(y_test, y_pred_GBC, rownames = ['Truth'], colnames =['Predicted'],
↪margins = True)
```

	precision	recall	f1-score	support
0	0.73	0.60	0.65	1095
1	0.73	0.83	0.78	1442

accuracy			0.73	2537
macro avg	0.73	0.71	0.71	2537
weighted avg	0.73	0.73	0.72	2537

```
[166]: Predicted    0     1   All
      Truth
      0         652   443 1095
      1         247  1195 1442
      All        899  1638 2537
```

```
[167]: #Applying K Nearest Neighbors
      KNN = KNeighborsClassifier().fit(X_train, y_train)

      y_pred_KNN = KNN.predict(X_test)

      print(classification_report(y_test, y_pred_KNN))

      #K Nearest Neighbors Confusion Matrix
      cm_LR = confusion_matrix(y_test,y_pred_KNN)
      pd.crosstab(y_test, y_pred_KNN, rownames = ['Truth'], colnames =['Predicted'],
      ↪margins = True)
```

		precision	recall	f1-score	support
	0	0.61	0.54	0.57	1095
	1	0.68	0.74	0.71	1442
	accuracy			0.65	2537
	macro avg	0.64	0.64	0.64	2537
	weighted avg	0.65	0.65	0.65	2537

```
[167]: Predicted    0     1   All
      Truth
      0         587   508 1095
      1         373  1069 1442
      All        960  1577 2537
```

```
[168]: #Applying Linear Discriminant Analysis
      LDA = LinearDiscriminantAnalysis().fit(X_train, y_train)

      y_pred_LDA = LDA.predict(X_test)

      print(classification_report(y_test, y_pred_LDA))

      #Linear Discriminant Analysis Confusion Matrix
      cm_LDA = confusion_matrix(y_test,y_pred_LDA)
```

```
pd.crosstab(y_test, y_pred_LDA, rownames = ['Truth'], colnames = ['Predicted'],  
↪margins = True)
```

	precision	recall	f1-score	support
0	0.66	0.57	0.61	1095
1	0.70	0.77	0.74	1442
accuracy			0.69	2537
macro avg	0.68	0.67	0.67	2537
weighted avg	0.68	0.69	0.68	2537

```
[168]: Predicted    0     1   All  
Truth  
0          622   473  1095  
1          325  1117  1442  
All         947  1590  2537
```

```
[169]: #Applying Gaussian Naive Bayes  
GNB = GaussianNB().fit(X_train, y_train)  
  
y_pred_GNB = GNB.predict(X_test)  
  
print(classification_report(y_test, y_pred_GNB))  
  
#Gaussian Naive Bayes Confusion Matrix  
cm_GNB = confusion_matrix(y_test, y_pred_GNB)  
pd.crosstab(y_test, y_pred_GNB, rownames = ['Truth'], colnames = ['Predicted'],  
↪margins = True)
```

	precision	recall	f1-score	support
0	0.56	0.62	0.59	1095
1	0.69	0.62	0.65	1442
accuracy			0.62	2537
macro avg	0.62	0.62	0.62	2537
weighted avg	0.63	0.62	0.63	2537

```
[169]: Predicted    0     1   All  
Truth  
0          683   412  1095  
1          543   899  1442  
All        1226  1311  2537
```

```
[170]: #Applying Support Vector Machine
SVM = SVC(kernel="rbf", random_state=None, probability=True, cache_size=500,
↳gamma=0.1).fit(X_train, y_train)

y_pred_SVM = SVM.predict(X_test)

print(classification_report(y_test, y_pred_SVM))

#Support Vector Machine Confusion Matrix
cm_SVM = confusion_matrix(y_test,y_pred_SVM)
pd.crosstab(y_test, y_pred_SVM, rownames = ['Truth'], colnames = ['Predicted'],
↳margins = True)
```

	precision	recall	f1-score	support
0	0.75	0.67	0.71	1095
1	0.77	0.83	0.80	1442
accuracy			0.76	2537
macro avg	0.76	0.75	0.75	2537
weighted avg	0.76	0.76	0.76	2537

```
[170]: Predicted    0    1   All
Truth
0          730   365 1095
1          242 1200 1442
All         972 1565 2537
```